



Compress attention. Never compress evidence.

A role-aware longitudinal care note where every high-priority statement carries its reason, trust state, and exact immutable source.

PRODUCT THESIS

The glance is an attention budget, not a second timeline.

Three bounded lanes - Act now, Watch, Awaiting - expose the next action while one click restores the full evidence trail.

1.425 ms

MEASURED WARM-PATH P95


159 tests

BACKEND AND REQUIRED MICRO-TESTS

100%

LINE + BRANCH COVERAGE

ARCHITECTURE



TRUST CONTRACT

- Exact-source highlights bind an entry, immutable version, offsets, quote, source URI, and SHA-256 hash.
- AI content starts as proposed; human-authored, staff-verified, clinician-confirmed, and superseded are visible states.
- Citation-first review returns only authorized claims, exact quotes, actions, and conflicts; unsupported questions receive an explicit abstention.
- Server-side clinic and role policies protect every object. Patient responses are built from an allow-list projection.
- Revert appends a version. Same-section stale writes return 409; independent role-owned sections do not overwrite.
- The Delta Lens reports temporal change and uncertainty, but refuses a causal claim without an estimand and design.

VERIFICATION CONTRACT

4 x 100%

FRONTEND COVERAGE + 7 E2E

900

PROPERTY EXAMPLES

0 known

DEPENDENCY ADVISORIES

0


MUTATION SURVIVORS



A shared narrative with immutable ownership.

Full snapshots favor auditability, deterministic revert, and stable provenance in the prototype. Production can add delta compression behind the same logical model.

COMPREHENSIVE DATA SCHEMA



AUTHORIZATION MATRIX

Object/action	Patient	Staff	Clinician	Admin
Patient-safe summary	Read	Read	Read/write	Read
Staff section	No access	Read/write	Read only	Read only
Clinician section	No access	Read only	Read/write	Read only
Raw AI + comments	No access	Read/review	Read/review	Read
Audit + retention	No access	No access	No access	Operate

ATOMIC INTEGRITY PATH

- 1. Resolve actor identity from the server-side membership; ignore client role claims.
- 2. Resolve the object inside clinic scope and check role ownership and visibility.
- 3. Validate expected version. A stale same-section mutation fails with a deterministic conflict receipt.
- 4. Append the snapshot, move the current pointer, append metadata-only audit, and refresh the glance projection.
- 5. Serialize only fields allowed for the actor; provenance resolution repeats scope and integrity checks.



Learning earns influence under safety constraints.

The adaptive system changes ranking, not clinical facts. New policies stay in shadow mode until support, uncertainty, and governance are credible.

TRANSPARENT IMPORTANCE

base = risk + action + safety + recency + pin
adaptive = clip(pooled posterior shift, -0.75, +0.75)
rank = hard safety band, then base + adaptive

- Role and clinic Beta posteriors shrink sparse feedback.
- Critical, allergy, medication, and urgent work are protected.
- Accept, reject, and pin remain one-action human controls.

SHADOW OFF-POLICY EVALUATION

No estimate yet


The seeded build correctly reports insufficient interaction data instead of manufacturing uplift.

REQUIRED DIAGNOSTICS

Propensity overlap - effective sample size - DR uncertainty - assumptions - policy version

Ranking proxy only - never a patient-outcome claim

PRIVACY AND RETENTION



Decay applies to recomputable derived caches. Immutable versions, audit metadata, open tasks, conflicts, pins, allergies, medication safety, and active critical evidence remain protected.

MEASURED EVIDENCE AND TRADE-OFFS

600 / 600

WARM READS SUCCEEDED

1.180 ms

LOCAL MEDIAN

1.425 ms

LOCAL P95

P95 <= 300 ms

BRIEF CONSTRAINT PASSED

50 warm-ups; loopback Uvicorn; seeded precomputed projection; no LLM on read path.

SELECTED PRIMARY SOURCES

- | | |
|--|---|
| [1] HL7 FHIR R5 Provenance | https://hl7.org/fhir/provenance.html |
| [2] W3C PROV-O | https://www.w3.org/TR/prov-o/ |
| [3] OWASP API1:2023 BOLA | https://owasp.org/API-Security/editions/2023/en/0xa1-broken-object-level-authorization/ |
| [4] NIST AI 600-1 Generative AI Profile | https://doi.org/10.6028/NIST.AI.600-1 |
| [5] Dai et al., AI-scribe safety signals, arXiv 2025 | https://arxiv.org/abs/2512.04118 |
| [6] Ambient scribe evaluation, JAMA Network Open, 2026 | https://jamanetwork.com/journals/jamanetworkopen/fullarticle/2843515 |
| [7] Mandym et al., CANDOR, CHIL 2026 | https://proceedings.mlr.press/v333/mandyam26a.html |